

Prefill

Process prompt once



Store each layer's
keys and values

Decode

One new query per step



Read old K/V
Append new K/V

At 256-token limit

Rebuild cropped context



Old states are invalid
for exact crop parity

Cache state belongs to one request; training has no KV cache.

MHA: 6 MiB GQA candidate: 1.5 MiB (FP32, one full request)